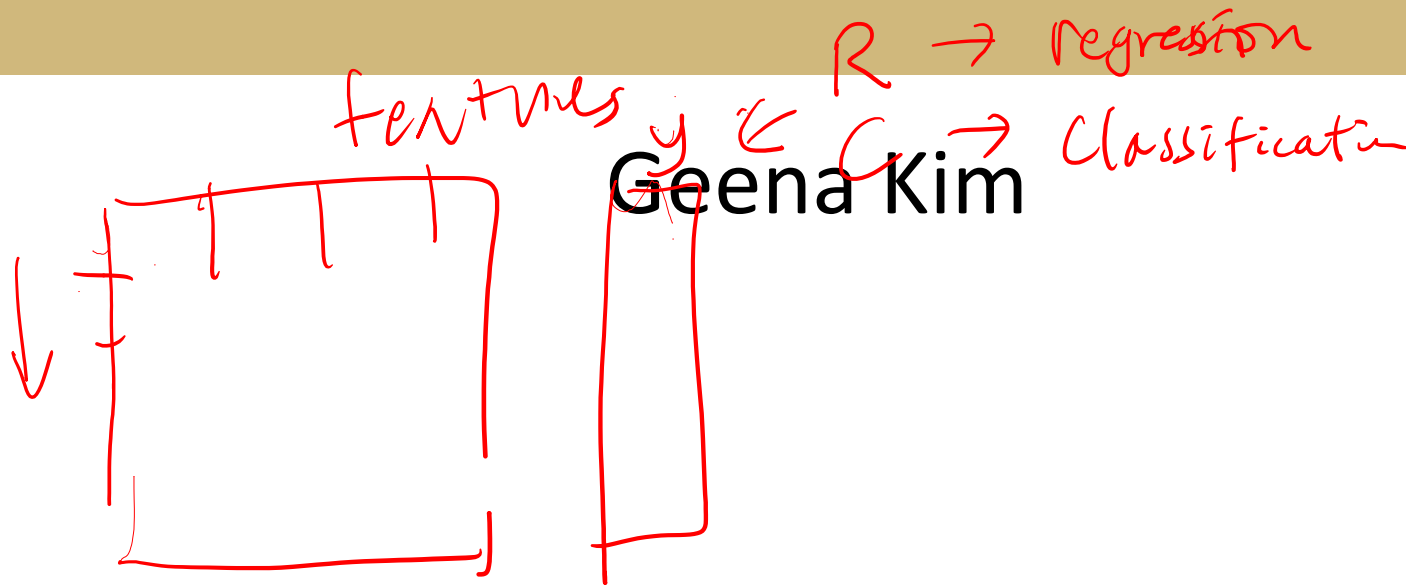


Linear Regression




Example

price	bedrooms	bathrooms	sqft_living	sqft_lot	floors
221900	3	1.00	1180	5650	1.0
538000	3	2.25	2570	7242	2.0
180000	2	1.00	770	10000	1.0
604000	4	3.00	1960	5000	1.0
510000	3	2.00	1680	8080	1.0


Data columns (total 21 columns):

id	21613	non-null	int64
date	21613	non-null	object
price	21613	non-null	int64
bedrooms	21613	non-null	int64
bathrooms	21613	non-null	float64
sqft_living	21613	non-null	int64
sqft_lot	21613	non-null	int64
floors	21613	non-null	float64
waterfront	21613	non-null	int64
view	21613	non-null	int64
condition	21613	non-null	int64
grade	21613	non-null	int64
sqft_above	21613	non-null	int64
sqft_basement	21613	non-null	int64
yr_built	21613	non-null	int64
yr_renovated	21613	non-null	int64
zipcode	21613	non-null	int64
lat	21613	non-null	float64
long	21613	non-null	float64
sqft_living15	21613	non-null	int64
sqft_lot15	21613	non-null	int64


Example



$f(x)$



Example



Simple Linear Regression

Intercept Slope Residual

$$\underline{Y} = \underline{\beta_0} + \underline{\beta_1} \underline{X} + \underline{\epsilon}$$

Coefficients, or Parameters

$$y_i - \hat{y}_i = \epsilon_i$$

Model


$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$$

Estimated values

Example

```
import statsmodels.formula.api as smf
model = smf.ols(formula='price~sqft_living', data=sub).fit()
model.summary()
```


Washington Home Sales Price 2014-2015



OLS Regression Results

Dep. Variable:	price	R-squared:	0.436			
Model:	OLS	Adj. R-squared:	0.436			
Method:	Least Squares	F-statistic:	1.651e+04			
Date:	Tue, 14 Jan 2020	Prob (F-statistic):	0.00			
Time:	19:04:13	Log-Likelihood:	-2.9523e+05			
No. Observations:	21402	AIC:	5.905e+05			
Df Residuals:	21400	BIC:	5.905e+05			
Df Model:	1					
Covariance Type:	nonrobust					
			95% CI			
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	1.24e+04	4305.280	2.881	0.004	3964.466	2.08e+04
sqft_living	251.0967	1.954	128.505	0.000	247.267	254.927

R-squared



$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$\text{TSS} = \sum_{i=1}^n (y_i - \bar{y})^2$$

$$\text{R}^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$$

p-value

Null hypothesis $H_0 : \beta_1 = 0$


Alternative hypothesis $H_A : \beta_1 \neq 0$

t-statistics

$$t = \frac{\hat{\beta}_1 - 0}{\text{SE}(\hat{\beta}_1)}$$


p-value

$$P > |t|$$



Accuracy of the coefficient estimates

```
import seaborn as sns
plt.figure(figsize=(10,6))
sns.regplot(x="sqft_living", y="price", data=x_sample, ci=95)
```



95% Confidence Interval

$$\left[\hat{\beta}_1 - 2 \cdot \text{SE}(\hat{\beta}_1), \hat{\beta}_1 + 2 \cdot \text{SE}(\hat{\beta}_1) \right]$$

Under the hood...

OLS Regression Results

Dep. Variable:	price	R-squared:	0.436
Model:	OLS	Adj. R-squared:	0.436
Method:	Least Squares	F-statistic:	1.651e+04
Date:	Tue, 14 Jan 2020	Prob (F-statistic):	0.00
Time:	19:04:13	Log-Likelihood:	-2.9523e+05
No. Observations:	21402	AIC:	5.905e+05
Df Residuals:	21400	BIC:	5.905e+05
Df Model:	1		
Covariance Type:	nonrobust		

Handwritten notes and equations illustrating the relationship between OLS and Maximum Likelihood Estimation (MLE).

MSE:
$$MSE = \frac{1}{2N} \sum_i^N (y_i - \hat{y}_i)^2$$

Log-Likelihood:
$$\prod_{i=1}^n p(y_i | x_i; b_0, b_1, s^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi s^2}} e^{-\frac{(y_i - (b_0 + b_1 x_i))^2}{2s^2}}$$

Maximum Likelihood:

$$L(b_0, b_1, s^2) = \log \prod_{i=1}^n p(y_i | x_i; b_0, b_1, s^2)$$


$$= \sum_{i=1}^n \log p(y_i | x_i; b_0, b_1, s^2)$$

$$= -\frac{n}{2} \log 2\pi - n \log s - \frac{1}{2s^2} \sum_{i=1}^n (y_i - (b_0 + b_1 x_i))^2$$

Gradient Descent

$$L = \text{MSE}$$

$$\text{MSE} = \frac{1}{2N} \sum_i^N (y_i - \hat{y}_i)^2$$



$$\frac{1}{2N} \sum_i^N (y_i - \beta_0 - \beta_1 x_i)^2$$







$$\beta_0 - \alpha \frac{\partial \text{MSE}}{\partial \beta_0}$$

Jacobian

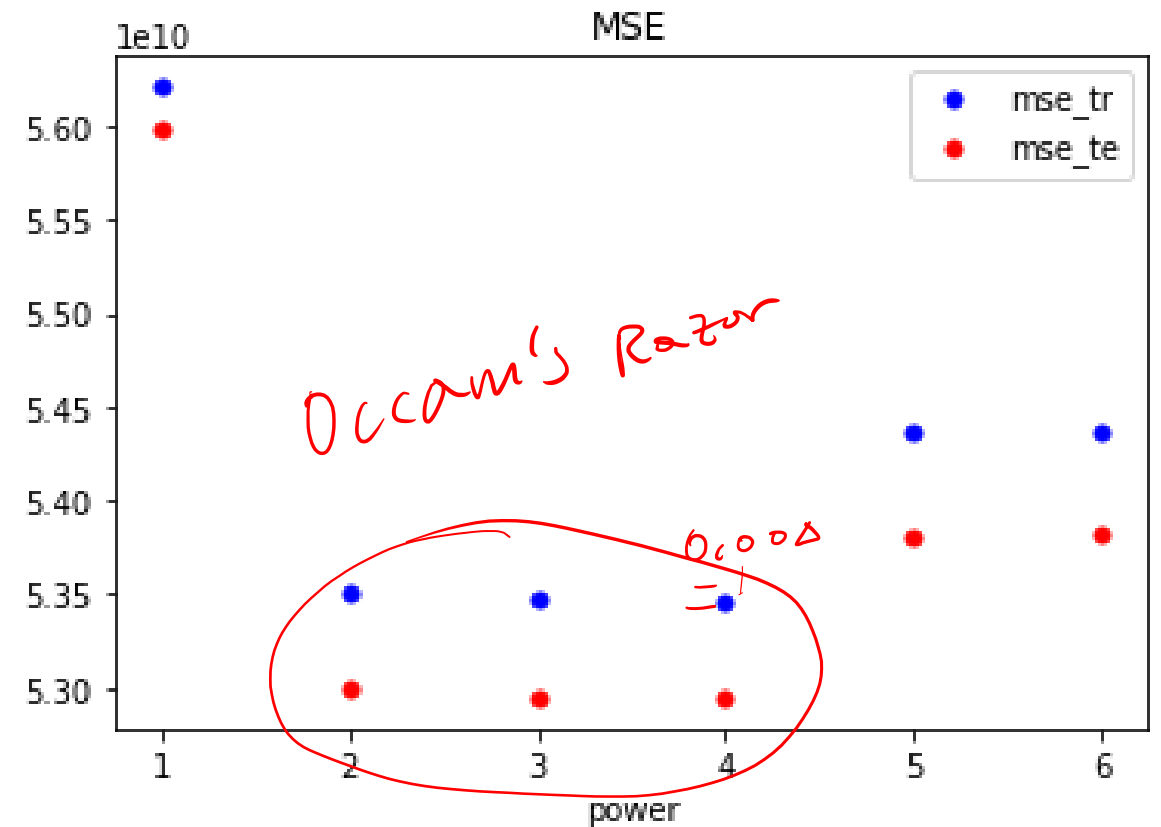
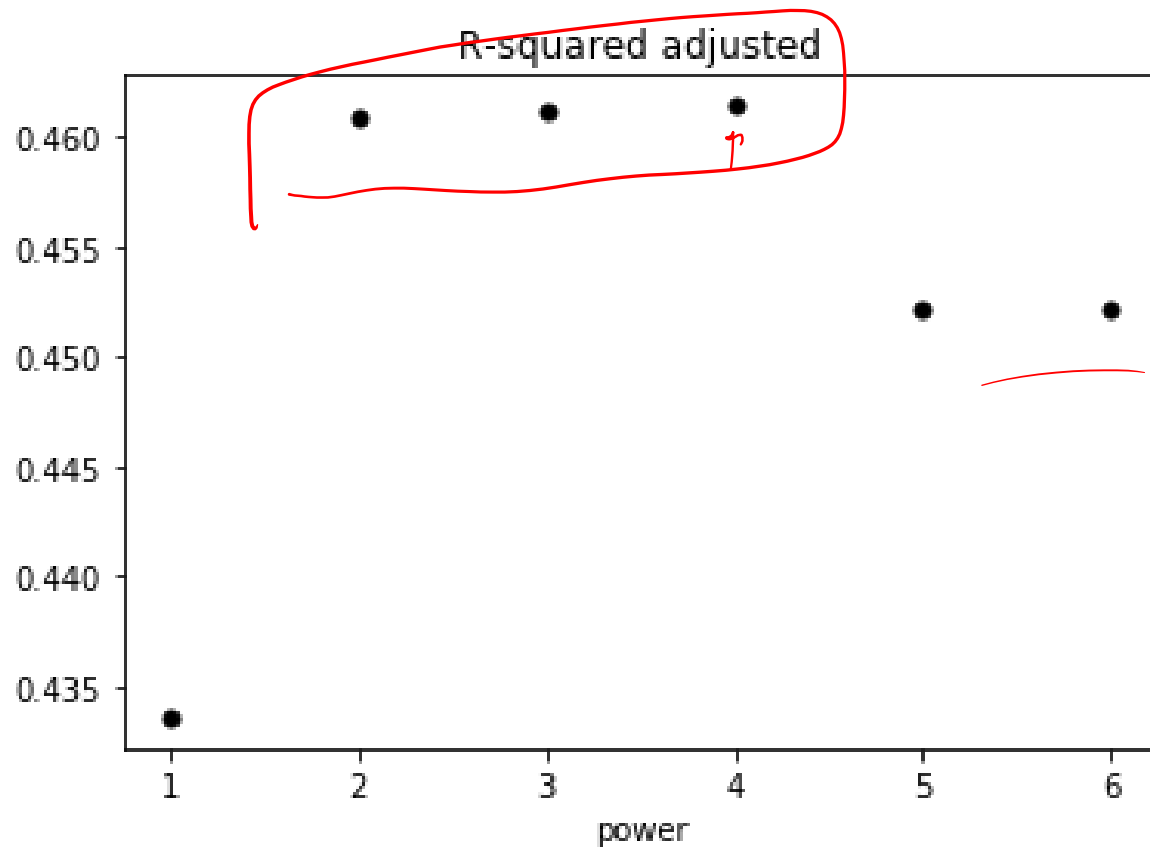
$$\nabla_{\beta}^2 \text{MSE}$$

Hessian


Polynomial Regression




Where to stop?



Bias-Variance Trade-off



high bias
low variance



low bias
high variance


Low Bias

High Bias


Low Variance

High Variance


a




b



c



d




Multi-linear Regression



Multilinear regression model

All predictors(variables) $X_1 \sim X_p$ are linear to Y

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$



Multilinear regression model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

β_j Average effect of X_j to Y when all other predictors fixed

Caution: In general,
predictors might be correlated

Collinearity

Caution: In general, predictors might be correlated

- When predictors are correlated, the coefficients estimates become inaccurate.
- Then it's called (multi-)collinear.
- The interpretation of the coefficient as the variable's contribution to the response becomes inaccurate.

Model selection

- Do I include all predictors or just a subset?

Forward Selection

step 1 $y = b_0$

step 2

{

```
p = []
for i = 1, ..., N
    y = b0 + bi * Xi
    p.append(p_value(y))
k = argmin(p)
y = b0 + bk * Xk
```

step 3 keep adding a predictor that gives the lowest p-value until it's satisfactory

Backward Selection

step 1 $y_0 = b_0 + b_1 * X_1 + b_2 * X_2 + \dots + b_N * X_N$

step 2

{

```
p = []
for i = 1, ..., N
    y = y0 - bi * Xi
    p.append(p_value(y))
k = argmax(p)
y = y0 - bk * Xk
```

step 3 keep removing a predictor that gives the largest p-value until it's satisfactory