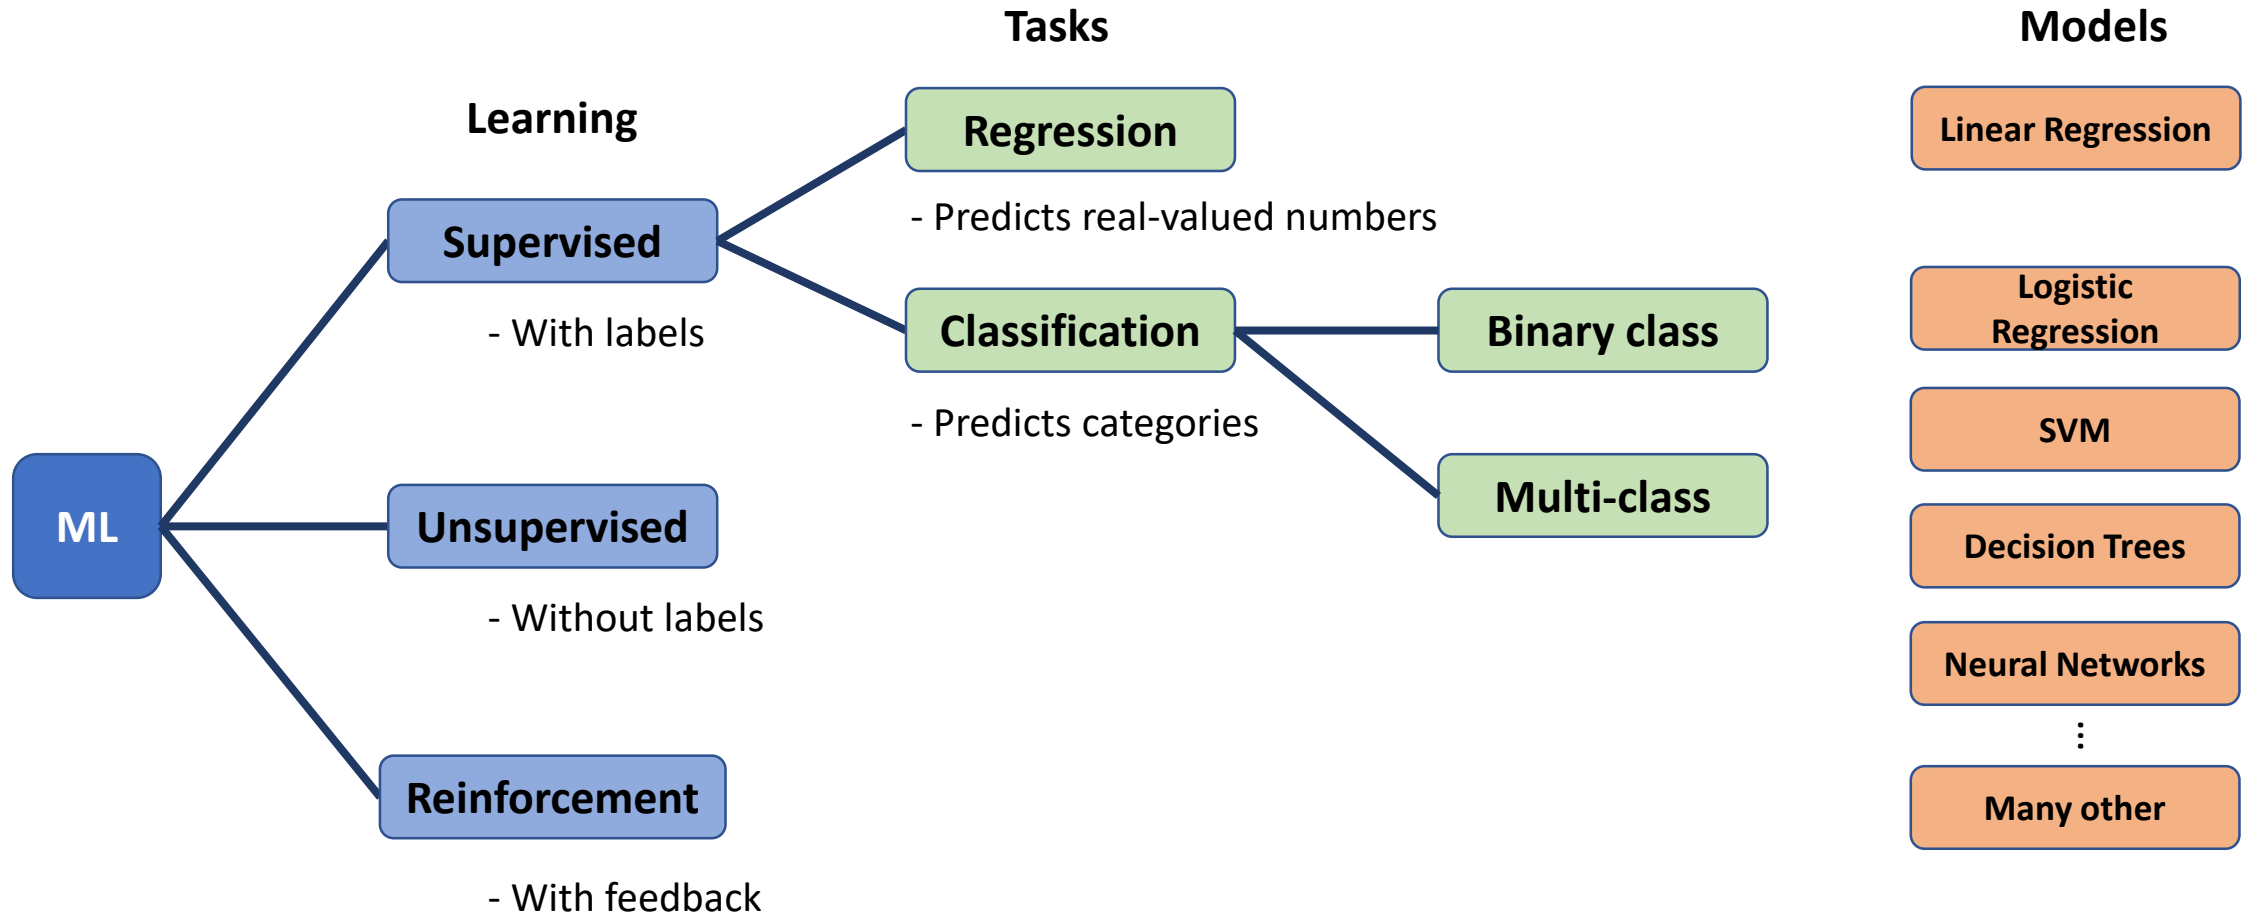


Logistic Regression

Geena Kim



Review- types of machine learning problems



Review- Linear Regression

$$\hat{y}^{(i)} = \mathbf{w} \cdot \mathbf{x}^{(i)} + b$$

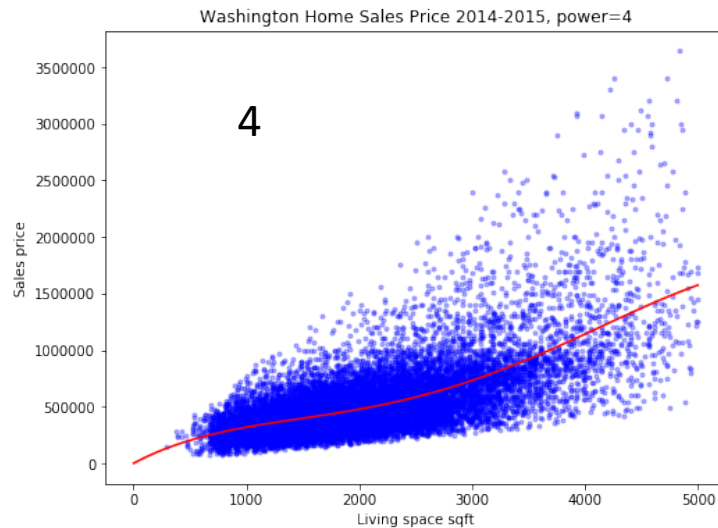
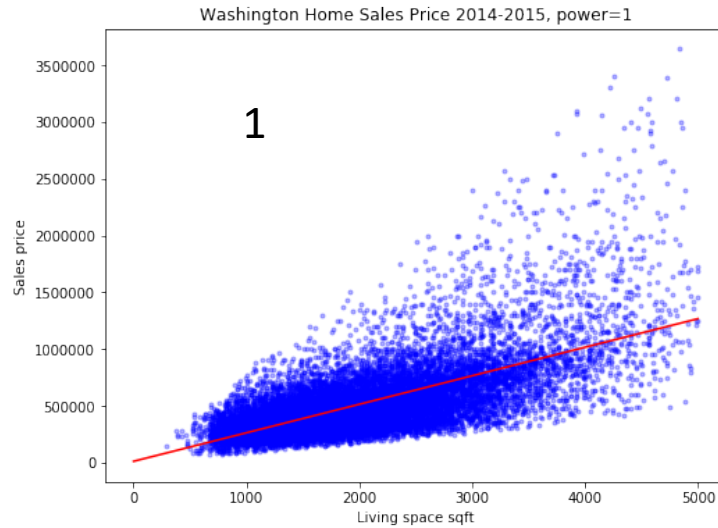
$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

$$\text{MSE} = \frac{1}{2m} \sum_i^m (\hat{y}^{(i)} - y^{(i)})^2$$

$$\hat{y}^{(i)} \in \mathbf{R}$$

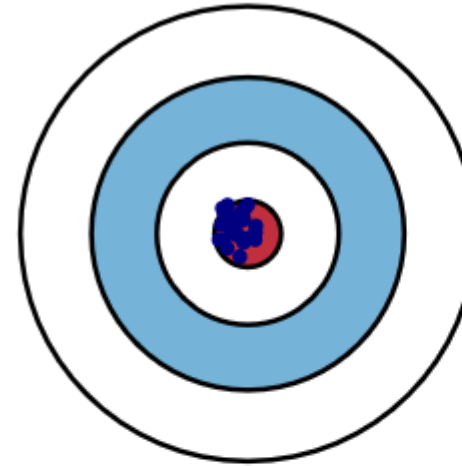


Bias-Variance Trade-off

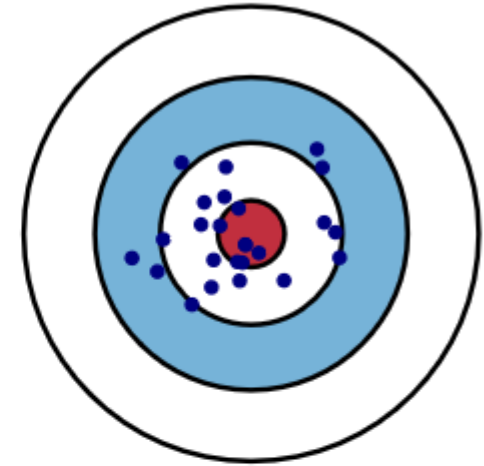


Low Bias

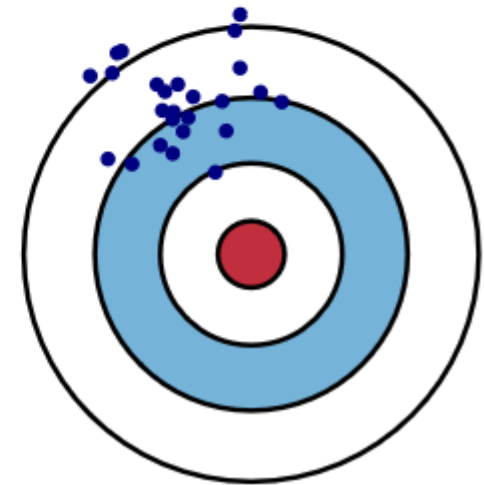
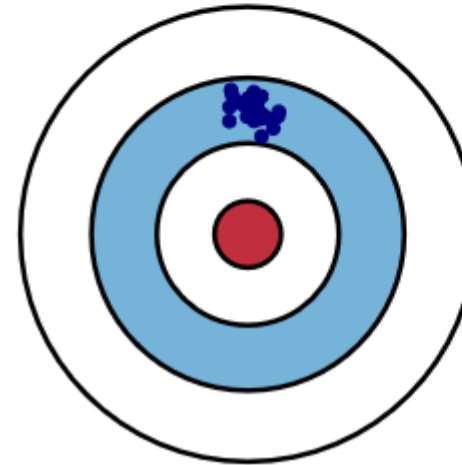
Low Variance



High Variance



High Bias

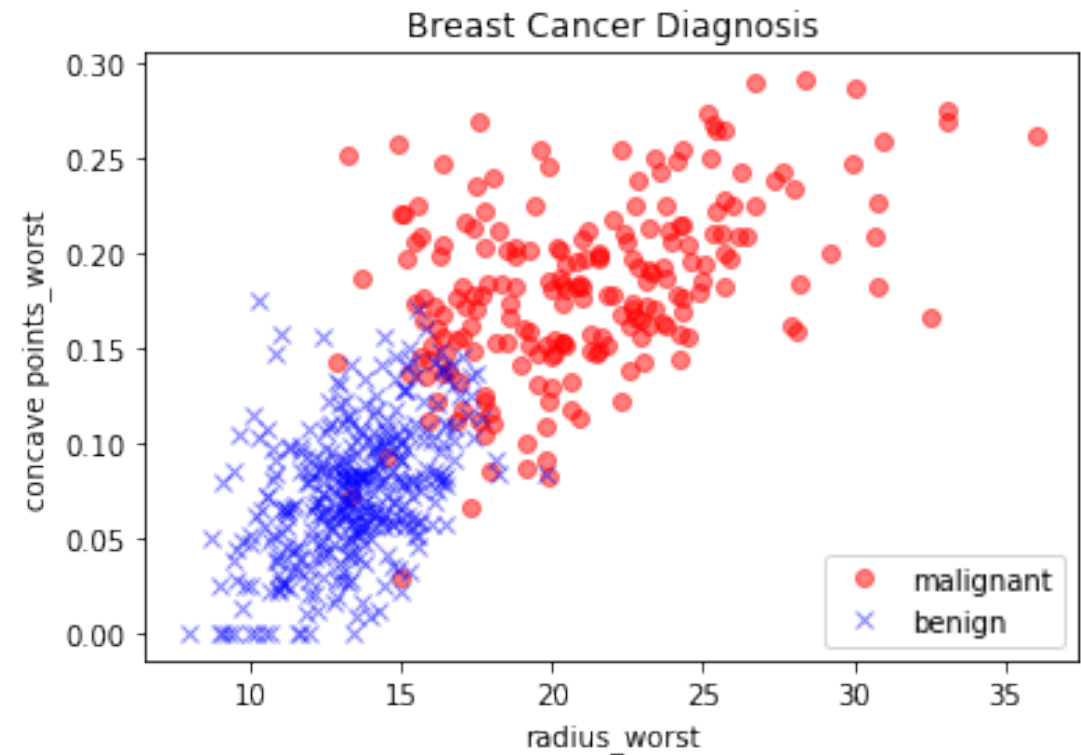



Binary Classification


Yes or No problem

- Creditcard Default
- Fraudulent Insurance Claim
- Spam Filtering
- Medical Diagnosis
- Survival Prediction
- Customer Retention
- Image Recognition

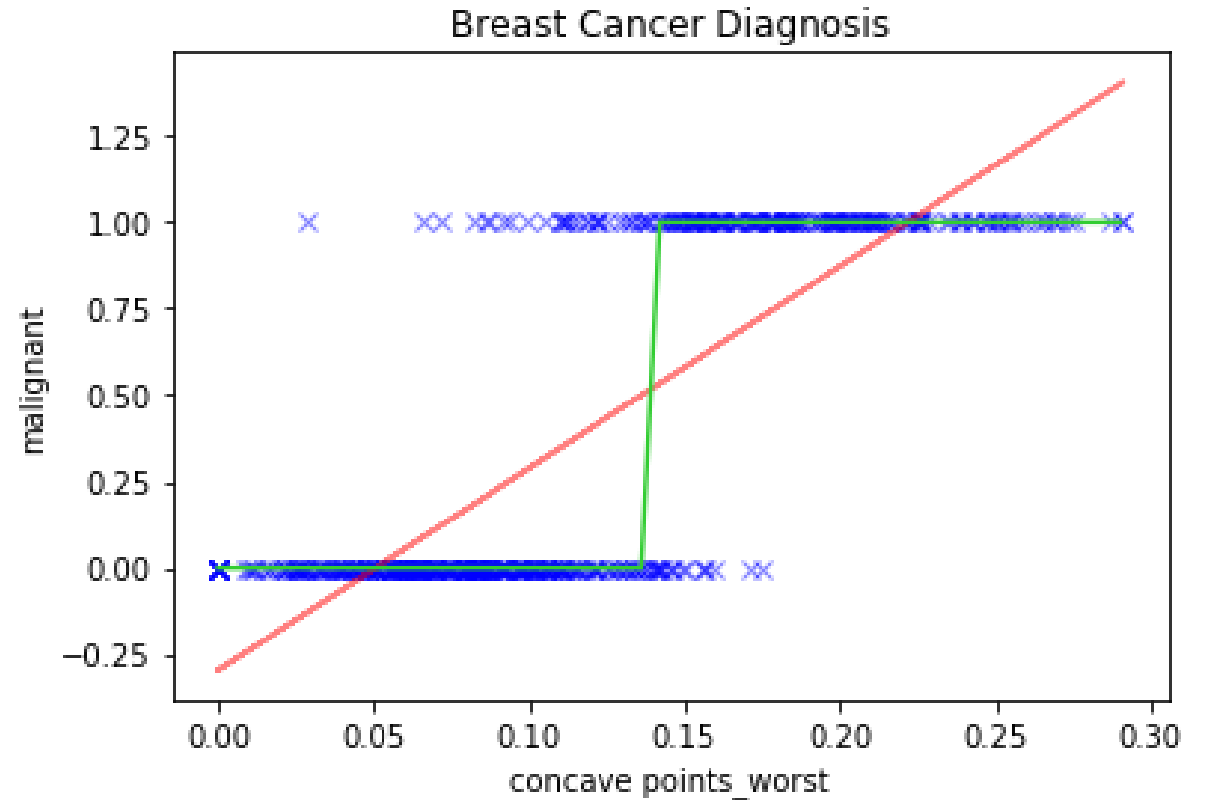
Binary Classification



Linear vs Logistic Regression




linreg.predict(x)




linreg.predict(x) > 0.5

Linear vs Logistic Regression

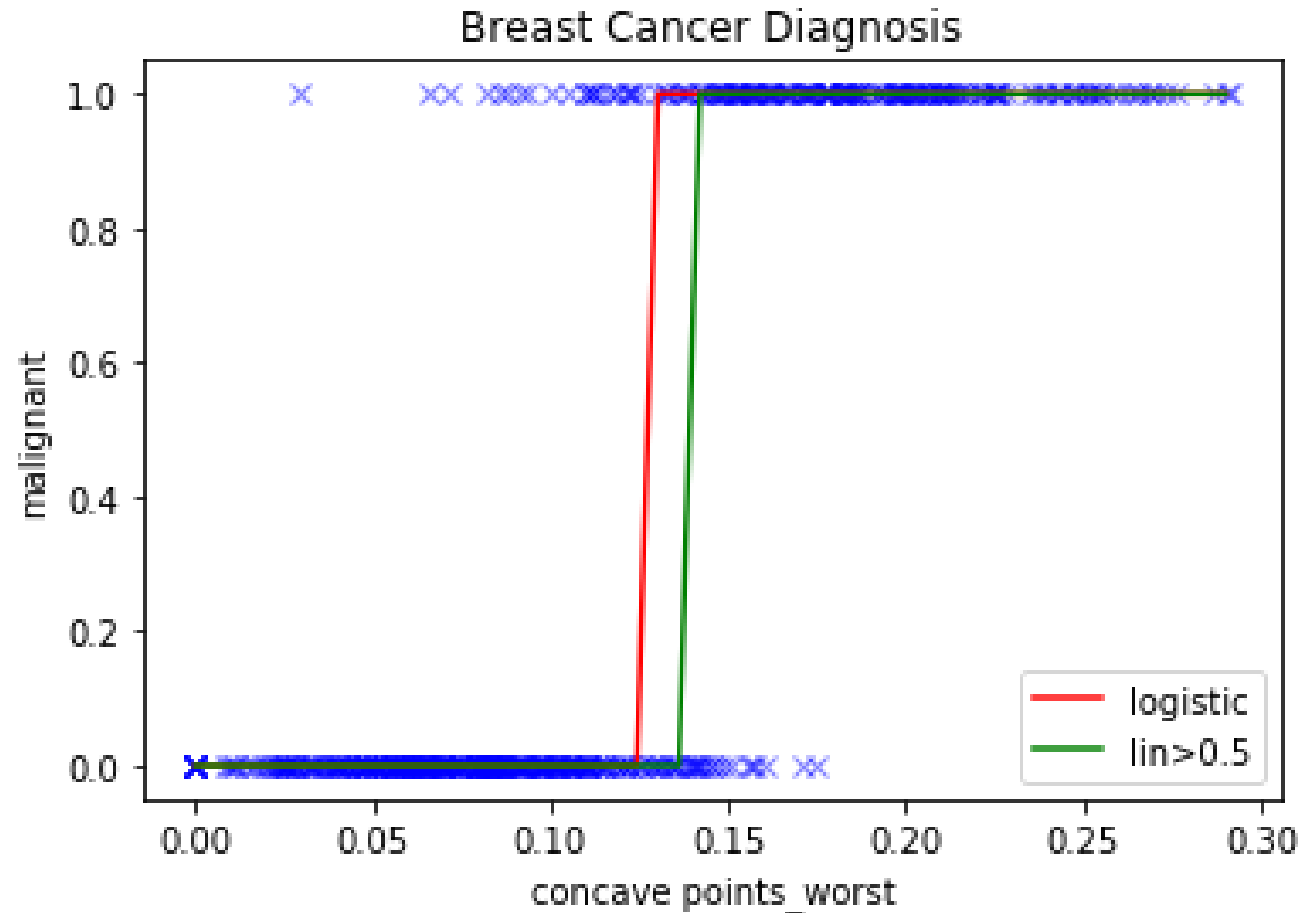


logreg.predict(x)



logreg.predict_proba(x)

Linear vs Logistic Regression

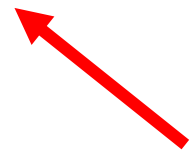


Logistic Function

$$P^{(i)} = \sigma(z^{(i)})$$


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z^{(i)} = \mathbf{W} \cdot \mathbf{X} + b$$



Called "logit" and is related to the decision boundary

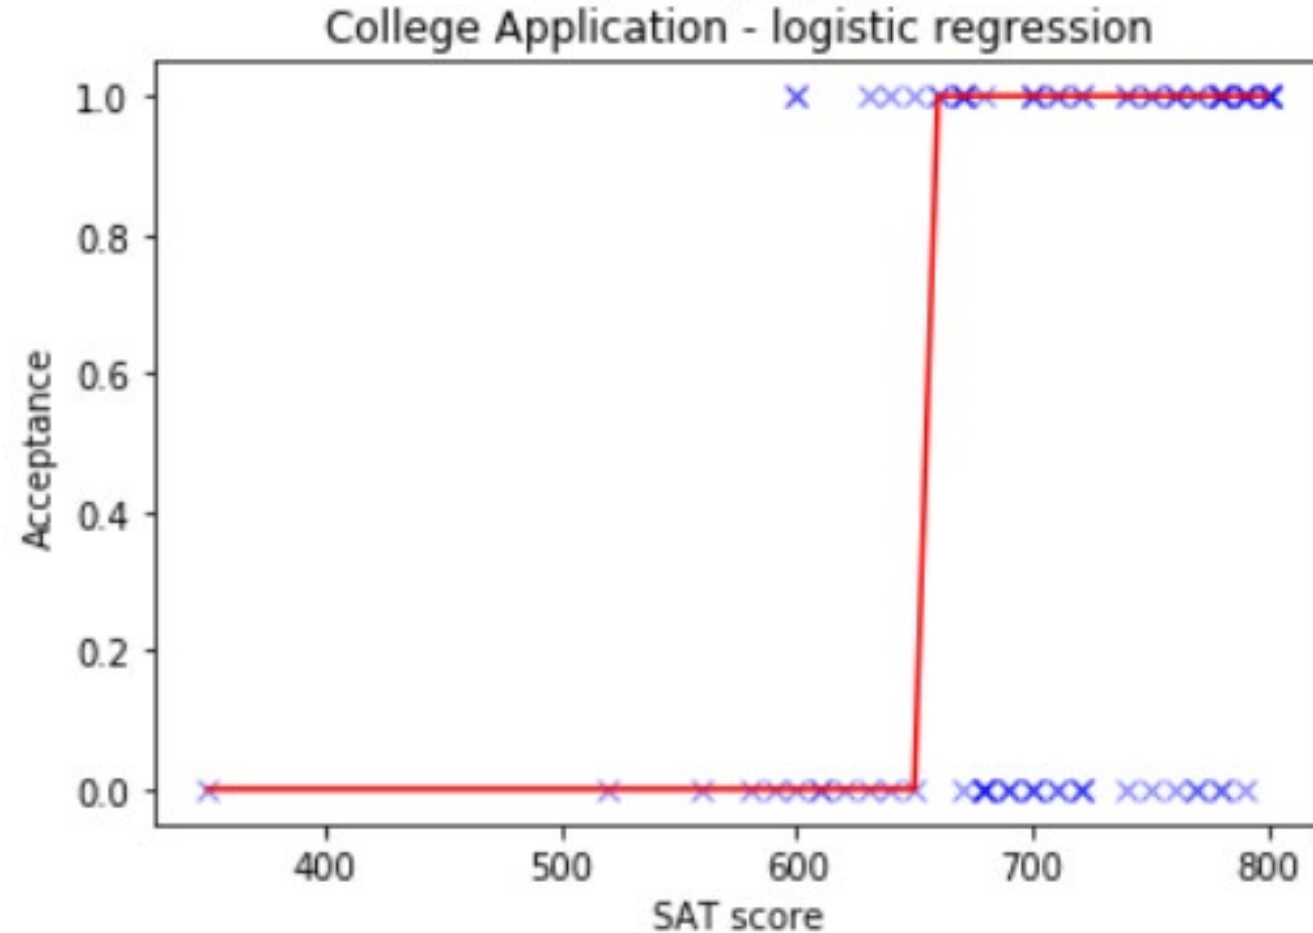
$$P^{(i)} \in \mathbb{R}[0, 1]$$



Logistic Regression- Univariate

University Acceptance


SAT_M	accept
690.0	0.0
710.0	1.0
790.0	1.0
770.0	0.0
770.0	1.0



Logistic Regression- Multivariate

Breast Cancer Diagnosis

radius_worst	concave points_worst	label
13.05	0.08263	0
16.39	0.16730	1
10.85	0.14650	0
21.86	0.15100	1
21.31	0.14900	1



$$z = 0.443 x_1 + 2.76 x_2 - 7.57 = 0$$

Estimating parameters in logistic regression

Maximum Likelihood

$$\ell(\beta_0, \beta) = \prod_{i:y_i=1} p(x_i) \prod_{i:y_i=0} (1 - p(x_i))$$

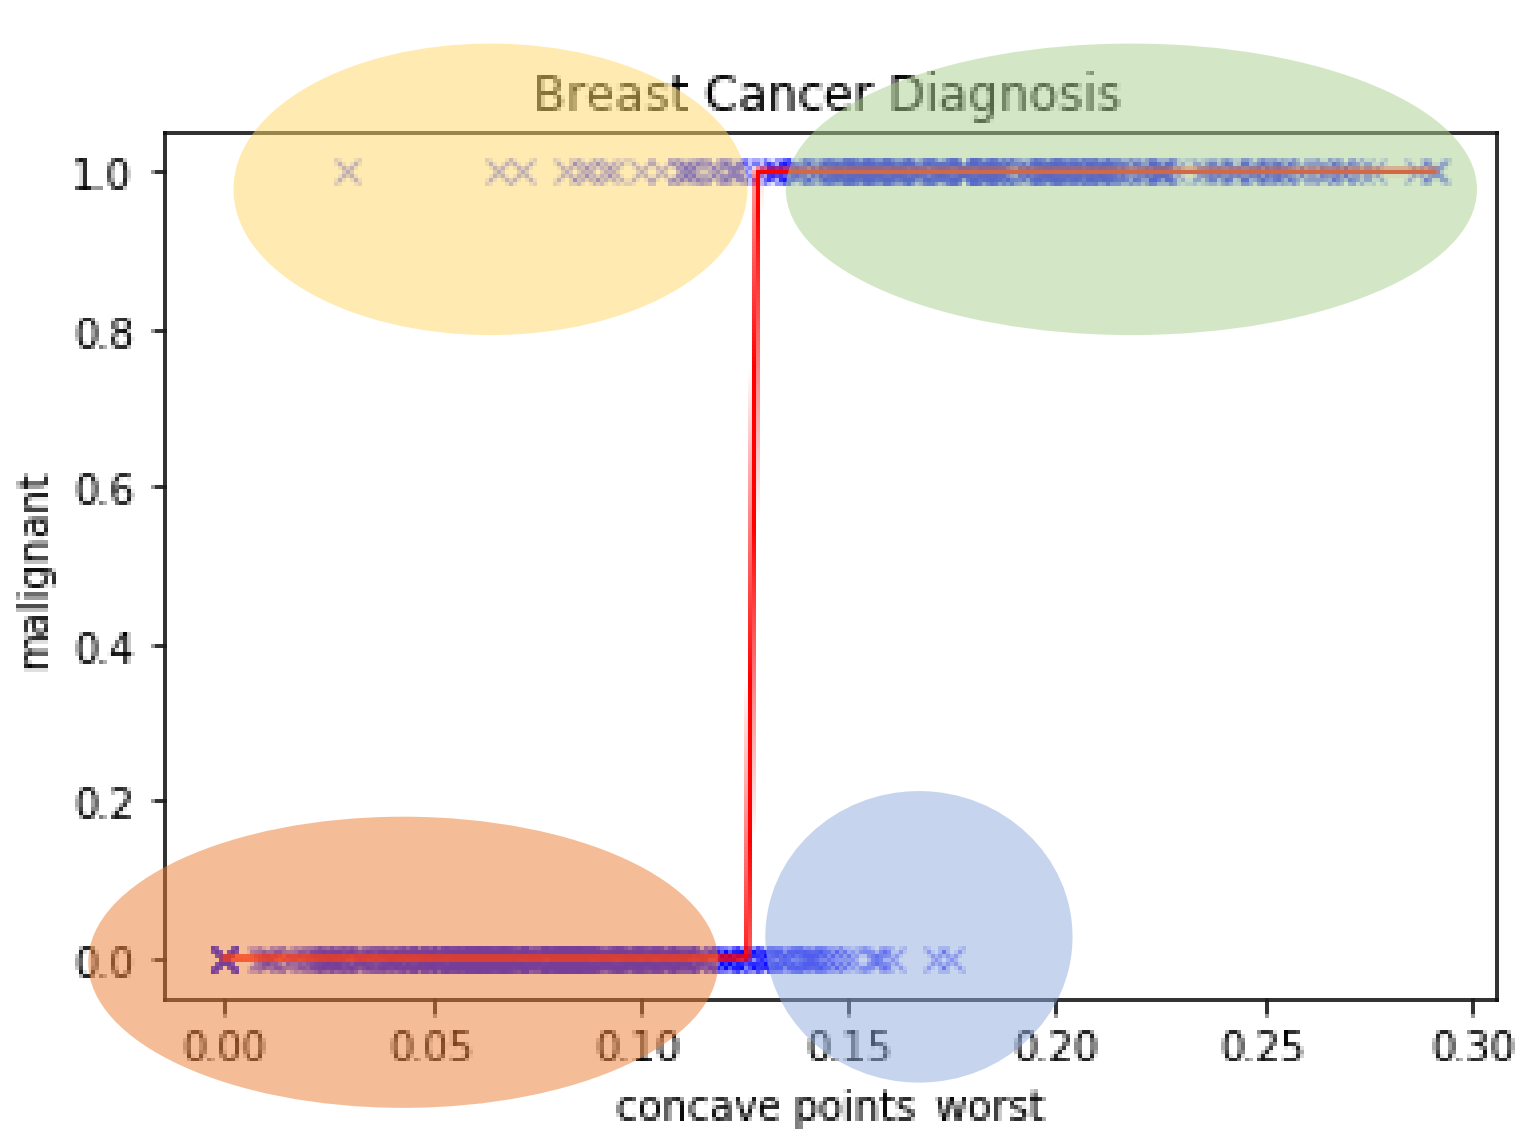
Estimating parameters in logistic regression

Cross Entropy

$$\mathcal{H}(P, Q) = - \sum_i P_i \log(Q_i)$$

$$= -\frac{1}{m} \sum_i^m y_i \log \hat{p}_i + (1 - y_i) \log (1 - \hat{p}_i)$$

Interpreting Logistic Regression Result



$Y_t = 1, Y_p = 1$
True Positive

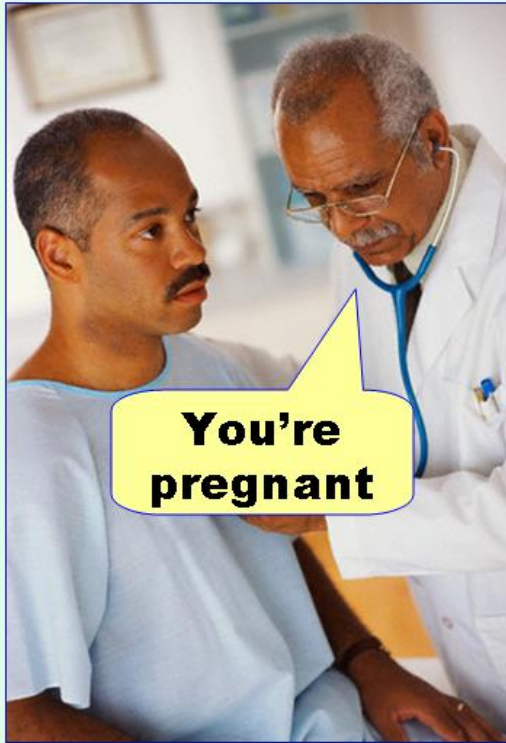
$Y_t = 0, Y_p = 0$
True Negative

$Y_t = 0, Y_p = 1$
False Positive


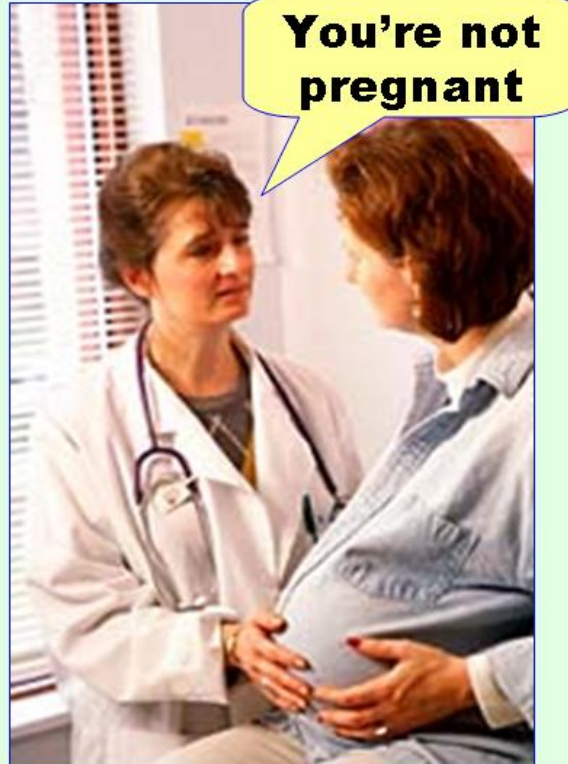
$Y_t = 1, Y_p = 0$
False Negative

Type I error and Type II error

Type I error
(false positive)



Type II error
(false negative)



Binary Classification Performance Metrics

Confusion Matrix


		Y _p	
		0	1
Y _t	0	70	1
	1	3	40

```
from sklearn.metrics import confusion_matrix  
confusion_matrix(y_true, y_pred)
```

```
pd.DataFrame(confusion_matrix(yt, yp, labels=[0,1]))
```

Performance Metrics- ROC, AUC

Receiver-Operating Characteristics Curve



Which Performance Metric should I choose?

- Accuracy
- Sensitivity, Recall, TPR
- Specificity, Sensitivity, TNR
- Precision, PPV
- False Positive Rate (fall-out)
- False Negative Rate (miss rate)
- F1 score
- AUC
- Confusion matrix

Loss: Why use Cross-Entropy, not Accuracy?

Cross Entropy

$$\mathcal{H}(P, Q) = - \sum_i P_i \log(Q_i) = - \frac{1}{m} \sum_i y_i \log \hat{p}_i + (1 - y_i) \log (1 - \hat{p}_i)$$

Accuracy

$$\frac{\text{TP} + \text{TN}}{\text{ALL}}$$

computed				targets					correct?

0.3	0.3	0.4		0	0	1	(democrat)		yes
0.3	0.4	0.3		0	1	0	(republican)		yes
0.1	0.2	0.7		1	0	0	(other)		no

computed				targets					correct?

0.1	0.2	0.7		0	0	1	(democrat)		yes
0.1	0.7	0.2		0	1	0	(republican)		yes
0.3	0.4	0.3		1	0	0	(other)		no

Scikit-Learn's logistic regression

[sklearn.linear_model](#).LogisticRegression

```
class sklearn.linear_model.LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True,  
intercept_scaling=1, class_weight=None, random_state=None, solver='warn', max_iter=100, multi_class='warn',  
verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

```
from sklearn.linear_model import LogisticRegression  
  
model = LogisticRegression().fit(X, y)
```

```
model.predict(X_test)
```

```
model.predict_proba(X_test)
```

https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html#sklearn.linear_model.LogisticRegression.decision_function

https://github.com/scikit-learn/scikit-learn/blob/1495f6924/sklearn/linear_model/logistic.py

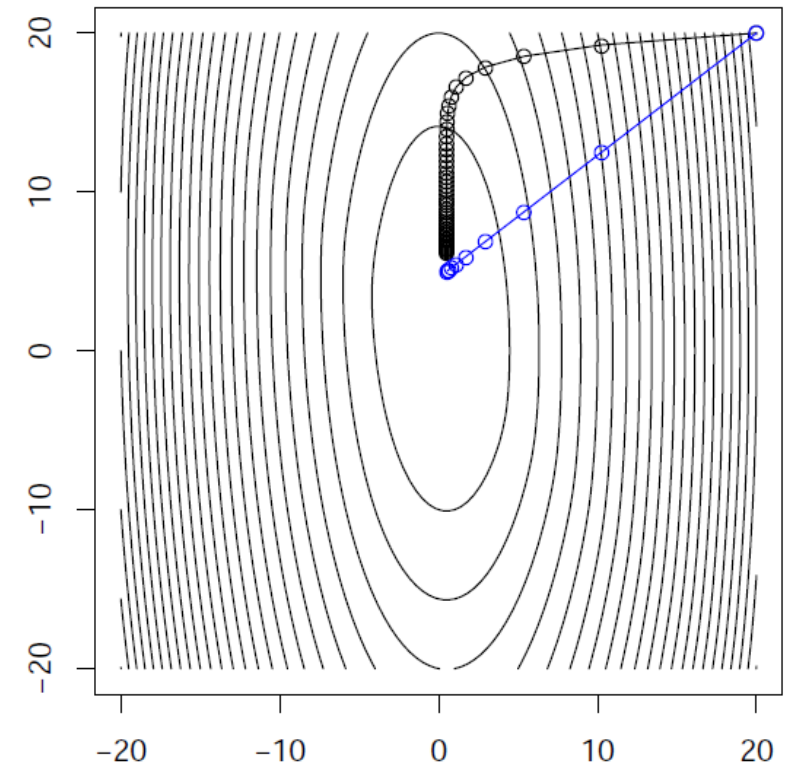
Under the Hood of sklearn's logistic regression

solver : str, {'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'}, optional
(default='liblinear').

liblinear (variant of Newton's method)

$$f(y) \approx f(x) + \nabla f(x)^T (y - x) + \frac{1}{2} (y - x)^T \nabla^2 f(x) (y - x)$$

$$x^+ = x - (\nabla^2 f(x))^{-1} \nabla f(x)$$



Using sklearn's LogisticRegression


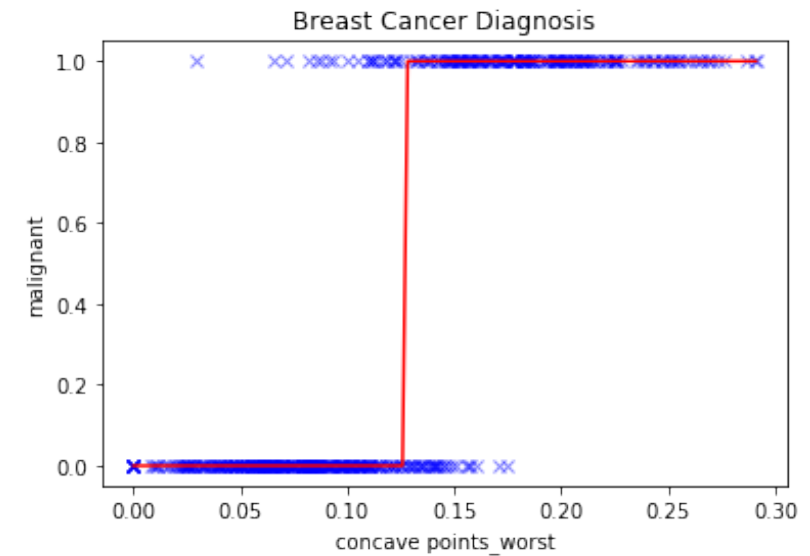
```
from sklearn.linear_model import LogisticRegression  
model = LogisticRegression().fit(X, y)
```

model.coef_

model.intercept_

Yp = model.predict(X_test)

P = model.predict_proba(X_test)



Using sklearn's LogisticRegression

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
from sklearn.linear_model import LogisticRegression as LR
```

```
clf = LR(class_weight="balanced", solver='liblinear').fit(X_train, y_train.ravel())
clf.score(X_test, y_test)
```

0.9649122807017544

```
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, recall_score
```

```
yp = clf.predict(X_test)
print('acc', accuracy_score(y_test, yp))
print('recall', recall_score(y_test, yp))
print('precision', precision_score(y_test, yp))
print('F1', f1_score(y_test, yp))
```

```
acc 0.9649122807017544
recall 0.9302325581395349
precision 0.975609756097561
F1 0.9523809523809524
```

```
pd.DataFrame(confusion_matrix(y_test, yp, labels=[0,1]))
```

	0	1
0	70	1
1	3	40

What about the statistics?

Another library

```
import statsmodels.api as sm
logit_model=sm.Logit(y_train,x_train)
result=logit_model.fit()
print(result.summary())
```

Optimization terminated successfully.
Current function value: 0.681033
Iterations 4

Logit Regression Results

Dep. Variable:	y	No. Observations:	455
Model:	Logit	Df Residuals:	454
Method:	MLE	Df Model:	0
Date:	Wed, 18 Sep 2019	Pseudo R-squ.:	-0.03232
Time:	19:23:16	Log-Likelihood:	-309.87
converged:	True	LL-Null:	-300.17
		LLR p-value:	nan

	coef	std err	z	P> z	[0.025	0.975]
x1	2.3970	0.731	3.279	0.001	0.964	3.830

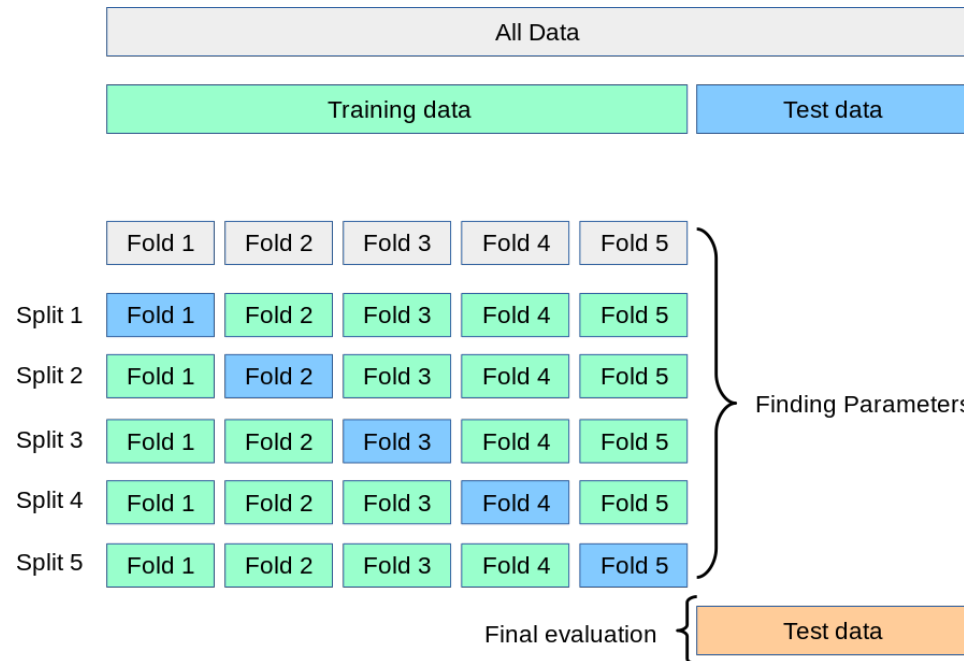
Bootstrap (Resample)

Next Lecture: Ways to train better

Regularization

```
class sklearn.linear_model.LogisticRegression(penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True,  
intercept_scaling=1, class_weight=None, random_state=None, solver='warn', max_iter=100, multi_class='warn',  
verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
```

Cross-Validation



```
class  
sklearn.linear_model.LogisticRegressionCV
```